

Celeriac: A Daikon .NET Front End

June 12, 2013

1 Introduction

Daikon is a tool that infers likely program invariants from program traces [1]. The Daikon tool relies on front ends, such as the Chicory front end for Java, to generate program traces to use for input. There are currently no available front ends for the .NET languages – C#, VB.NET, and F#.

The purpose of this document is to describe the architecture of Celeriac, a Daikon Front End for .NET. It will be updated throughout the development process as to accurately reflect the design and architecture of the system.

2 Design Goals

This section describes the design goals guiding the development of the .NET Daikon front end.

Work with any .NET program. While Celeriac will likely primarily be used with programs written in C#, we leverage the Common Language Runtime architecture to make the front end work with many .NET languages. Celeriac has been successfully used on F# and VB.NET in addition to C#. This compatibility is achieved by targeting the tool at binaries that have already been compiled to CIL, as opposed to targeting source code.

However, Celeriac must specially detect the presence of some language-specific constructs, to enable detection of meaningful properties. As an example of this F# uses its own list type instead of the .NET `System.Collections.List` type. To capture meaningful invariants over lists Celeriac treats a variable as a List if it has the `System.Collections.List` or the `FSharp.List` type.

Don't require modification of the developer's build process. Commercial build processes can become very complex over time. As a result, it may be unwieldy or even impossible to integrate new tools into the build chain. In order to promote tool use, Celeriac operates on the output of the build process: a managed binary executable or library. An added benefit of meeting this design goal is that Celeriac can run on programs even when the source is not available.

Don't modify files on disk. Modifying the contents of the disk to instrument the program, e.g., creating a patched version of the executable, is similarly undesirable. The creation or replacement of binary files opens the door for versioning issues such as accidentally running an instrumented

binary with incorrect options.

There are some situations meeting where this goal isn't possible. For example, when create an instrumented library that will be used by other programs the library must be saved to disk. Also, Celeriac is running on GUI programs using Windows Presentation Foundation (WPF), Celeriac must be run in off-line mode. This is because WPF programs load resources in a way that cannot be done in online mode, since the subject program is loaded into the custom launcher's `AppDomain`.

3 Tool Architecture

The front end has two major components: an IL rewriter that inserts instrumentation calls into procedures and a run-time system that, when called from the subject program, outputs variable values to construct a datatrace in the Daikon format.

3.1 IL Rewriter

The IL rewriter is implemented in `ILRewriter.cs` and rewrites .NET binaries by inserting calls to the Celeriac runtime library. It utilizes the CCIMetadata IL Rewriting library [2]. The CCIMetadata library provides an API for reading and writing CIL at a higher level than raw bytecode. The IL rewriter component of the application, named `ILRewriter`, was adapted from a CCIMetadata example program that printed the name of local variables. The rewriter loads a .NET Module in memory, walks over the IL and inserts instrumentation calls at the entrance and every exit of methods. The following figure illustrates the IL rewriter component.

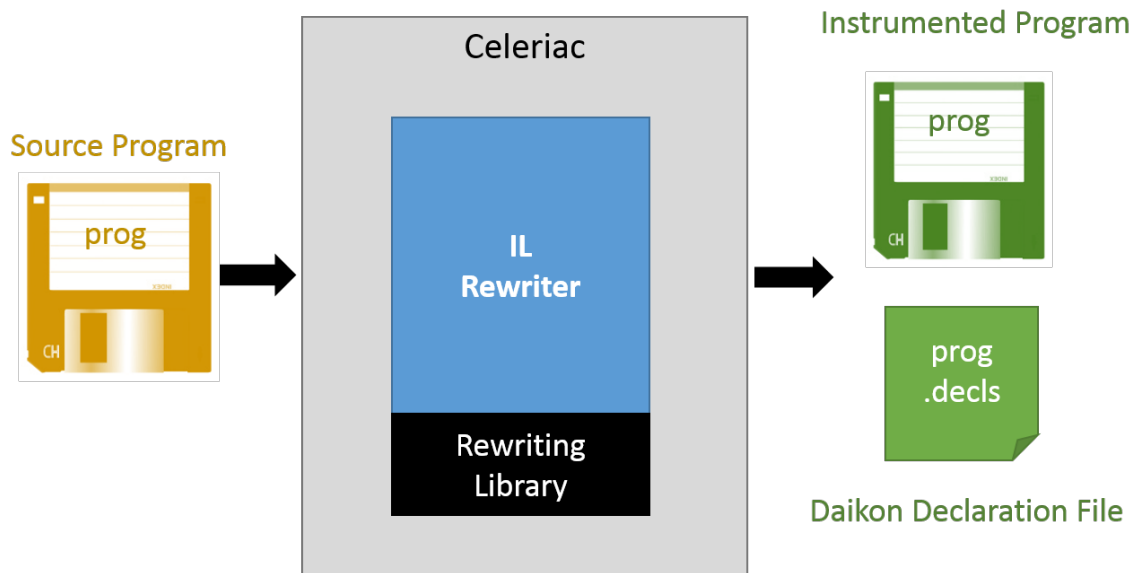


Figure 1: IL Rewriter Architecture

A Daikon datatrace file consists of two parts. The first, called the declaration component, is metadata describing the schema of each program point (method entrance or exit). While Celeriac is inserting instrumentation calls the IL rewriter also makes calls to the `DeclarationPrinter` class which writes the declaration portion of the datatrace file. The second component, the datatrace component, will be written by the program as it executes. When the IL rewriter is finished, it returns a memory stream which the front-end loads and executes, or saves to disk if desired.

3.2 Runtime Library

The runtime library creates the actual datatrace file that will be input to Daikon. The main method exposed to source programs prints the name of a variable, and its value, then prints the name and value of each of that variable's fields, pure methods, and elements if any. The following figure illustrates the functionality of the runtime library component.

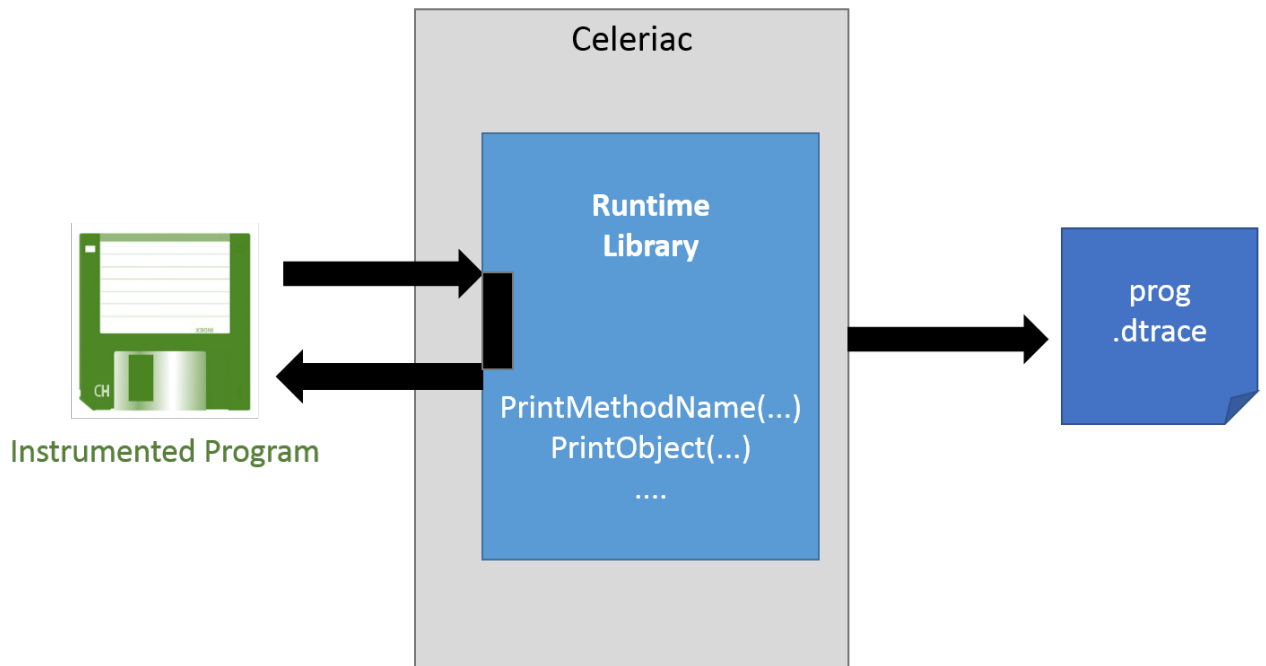


Figure 2: Runtime Library Architecture

The runtime library is implemented in the `VariableVisitor` class. When the `VisitVariable()` method is called from the instrumented program the value of the variable itself is printed. Then any static or instance fields of that variable, which are visited reflectively based on the declared type of the variable are also printed. The instrumentation code traverses data structures and enumerates variables that implement the `ILList`, `ISet`, or `Dictionary` interfaces, or their F# equivalents.

4 Design Decisions

The situation describes important design decisions, the alternatives, and the justification for the implemented choice.

4.1 Type Resolution

4.2 Instrumentation Trade-Off

4.3 Exceptional Post-Conditions

A Appendix: Project Schedule and Milestones

This section describes the project milestones and history.

A.1 10/2010 - 12/2010

Work began on the instrumentor this quarter. A first attempt was to use the unmanaged Profiler API. First C# reflection was used to analyze the program, and then the Profiler API's function enter/leave hooks were used to output the appropriate variable information. Manually editing the MSIL was extremely difficult and error-prone, especially for methods with multiple exit points and exception handling, so this approach was abandoned.

A pure IL rewriting approach using the CCIMetadata Library was attempted, where the output program was saved to disk as a new executable, however this violated the Disk Modifications Avoidance design goal. A hybrid of the previous two methods was considered where the CCIMetadata Library would create the new MSIL, and this would be handed to the unmanaged Profiler, to be substituted for the original method IL.

A presentation of the work to this point was presented to the UW CSE PL&SE group on December 13th.

A.2 1/2011 - 3/2011

The previously described method was almost immediately shown to be infeasible. The CCIMetadata rewritten code produced different method tokens than were present in the original program, and thus could not be used. The main developer on the CCIMetadata project confirmed this issue could not be resolved in context. Analysis of alternate methods for IL rewriting was completed. It was determined that using the program with instrumentation inserted by CCIMetadata, but loaded into memory and ran instead of written to disk was the best approach.

The remainder of the quarter was spend implementing the profiler and visitor, and getting them to conform to Daikon specification. At the conclusion of the quarter an alpha version of the front end was functional, and could generate simple datatrace files that Daikon could process.

A.3 4/2011 - 6/2011

During this quarter reflective visiting and array-style visiting of IEnumerable types were implemented. Most command-line arguments available to Chicory were documented. A test suite was added, modeled off the Kavasir (C++) front-end test suite. Chicory test sources were transliterated

from Java to C# and ran through the front-end. Output invariants were compared to the invariants produced by Chicory, and bugs in the front-end were fixed to rectify the differences. Exception handling was also added, a feature new to Daikon front-ends.

A.4 9/2011 - 12/2011

During this quarter the first code review of Celeriac was completed, resulting in many name improvements. Celeriac was extended to output multiple method exits points for methods that return more than once. A Google Code issue tracker was established. Celeriac was extended to work on programs with multiple threads.

A.5 1/2012 - 3/2012

During this quarter Celeriac was extended to run on pure methods and with the first additional language, F#. This introduced the complication that F# uses its own lists, and special handling was added for these.

A.6 4/2012 - 6/2012

During this quarter Celeriac's collection support was significantly expanded, specifically to include both F# and C# Sets and Maps. As well, support for generic types with multiple constraints was added as well. The documentation was significantly updated.

A.7 9/2012 - 12/2012

During this quarter the test suite was expanded and enhanced. WPF programs were tested for the first time, as were VB.NET programs. These each introduced bugs, which were fixed. Support was added for using Celeriac on programs exercised by external testing systems such as the Visual Studio testing system and NUnit.

A.8 1/2013 - 3/2013

During this quarter Celeriac was expanded to run on large sample .NET programs, for the research paper. This included running Celeriac on programs in a variety of languages and types, e.g. libraries, GUI applications, etc. A new command line option for WPF programs was added. Bugs were fixed with elimination of jagged and multi-dimensional arrays and type issues. The `-std-visibility` flag was significantly broken before, and revamped. Pure method instrumentation was improved. Celeriac was run on the Mono test suite, giving higher confidence in its robustness.

A.9 4/2013 - 6/2013

In this quarter much of the work was focused on wrapping up Celeriac, rather than implementation of new features. The first full release of Celeriac was prepared, and made on June 14. Developer and user documentation were updated and expanded. Many old issues were fixed for the release. Remaining issues were examined and received detailed explanations for new developers. A presentation on Celeriac was given to the PLSE group on May 17th.

References

- [1] *The Daikon Invariant Detector User Manual*, September 2010. <http://groups.csail.mit.edu/pag/daikon/download/doc/daikon.html>.
- [2] Microsoft Research. <https://ccimetadata.codeplex.com/>.